

# Developer Operations

---

AWS Services

 <b>Compute</b>	 <b>Developer Tools</b>	 <b>Analytics</b>	 <b>Application Services</b>
EC2	CodeStar	Athena	Step Functions
Elastic Container Service	CodeCommit	EMR	SWF
Lightsail	CodeBuild	CloudSearch	API Gateway
Elastic Beanstalk	CodeDeploy	Elasticsearch Service	Elastic Transcoder
Lambda	CodePipeline	Kinesis	
Batch	X-Ray	Data Pipeline	
 <b>Storage</b>	 <b>Management Tools</b>	 <b>Artificial Intelligence</b>	 <b>Messaging</b>
S3	CloudWatch	Lex	Simple Queue Service
EFS	CloudFormation	Amazon Polly	Simple Notification Service
Glacier	CloudTrail	Rekognition	Simple Email Service
Storage Gateway	Config	Machine Learning	
 <b>Database</b>	OpsWorks	 <b>Internet Of Things</b>	 <b>Business Productivity</b>
RDS	Service Catalog	AWS IoT	WorkDocs
DynamoDB	Trusted Advisor	AWS Greengrass	WorkMail
ElastiCache	Managed Services		Amazon Chime
Amazon Redshift	 <b>Security, Identity &amp; Compliance</b>	 <b>Contact Center</b>	 <b>Desktop &amp; App Streaming</b>
 <b>Networking &amp; Content Delivery</b>	IAM	Amazon Connect	WorkSpaces
VPC	Inspector		AppStream 2.0
CloudFront	Certificate Manager	 <b>Game Development</b>	
Direct Connect	Directory Service	Amazon GameLift	
Route 53	WAF & Shield		
 <b>Migration</b>	Artifact	 <b>Mobile Services</b>	
AWS Migration Hub	Amazon Macie	Mobile Hub	
Application Discovery Service	CloudHSM	Cognito	
Database Migration Service		Device Farm	
Server Migration Service		Mobile Analytics	
Snowball		Pinpoint	

**Mar. 2019 –  
100+ Services**

# AWS Shared Responsibility Model

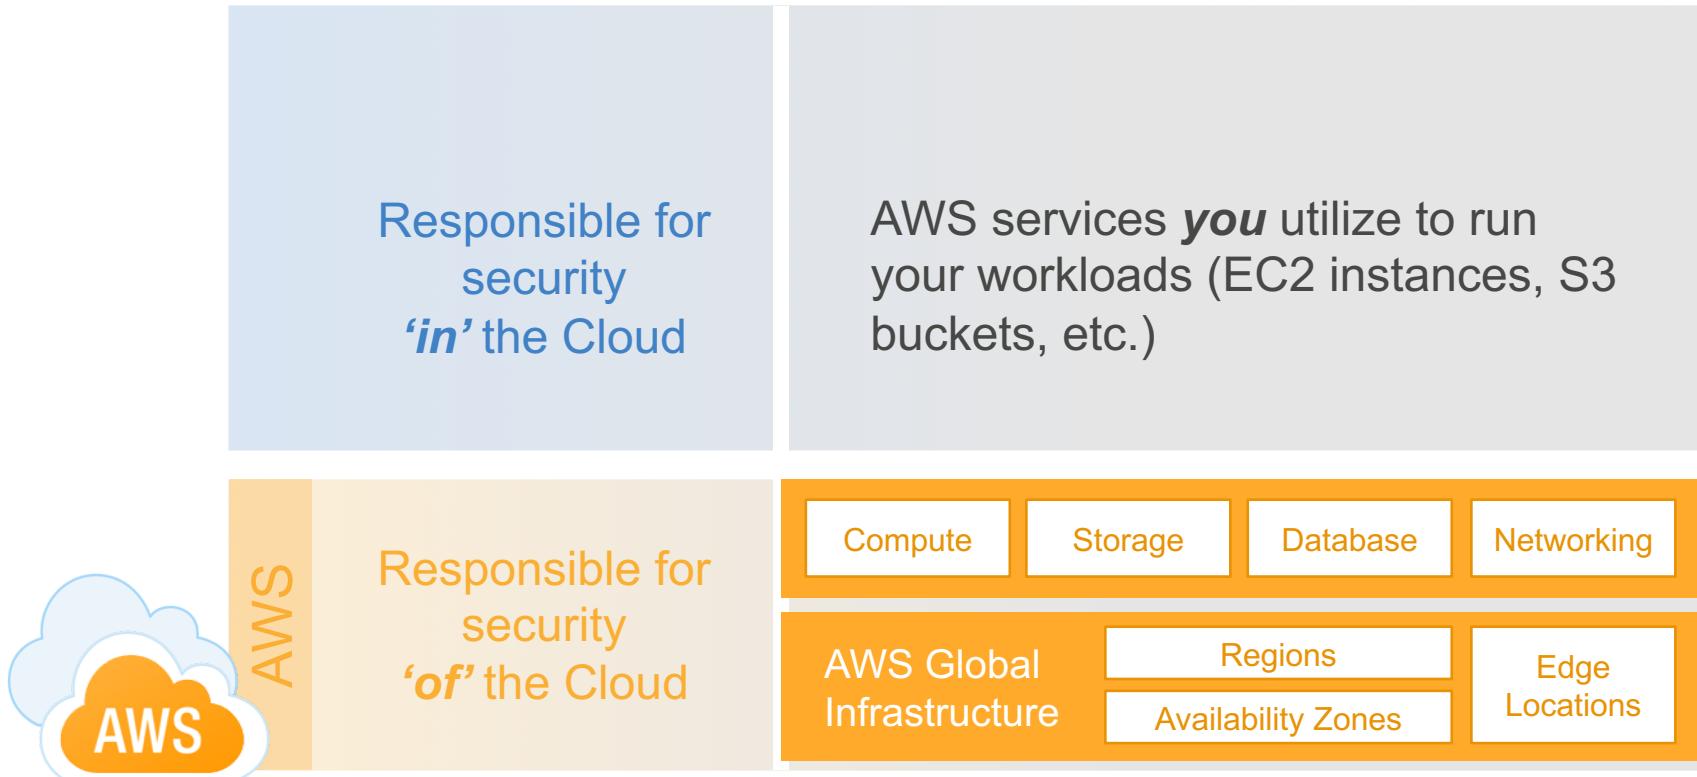
Responsible for  
security  
**'in'** the Cloud

AWS services ***you*** utilize to run  
your workloads (EC2 instances, S3  
buckets, etc.)

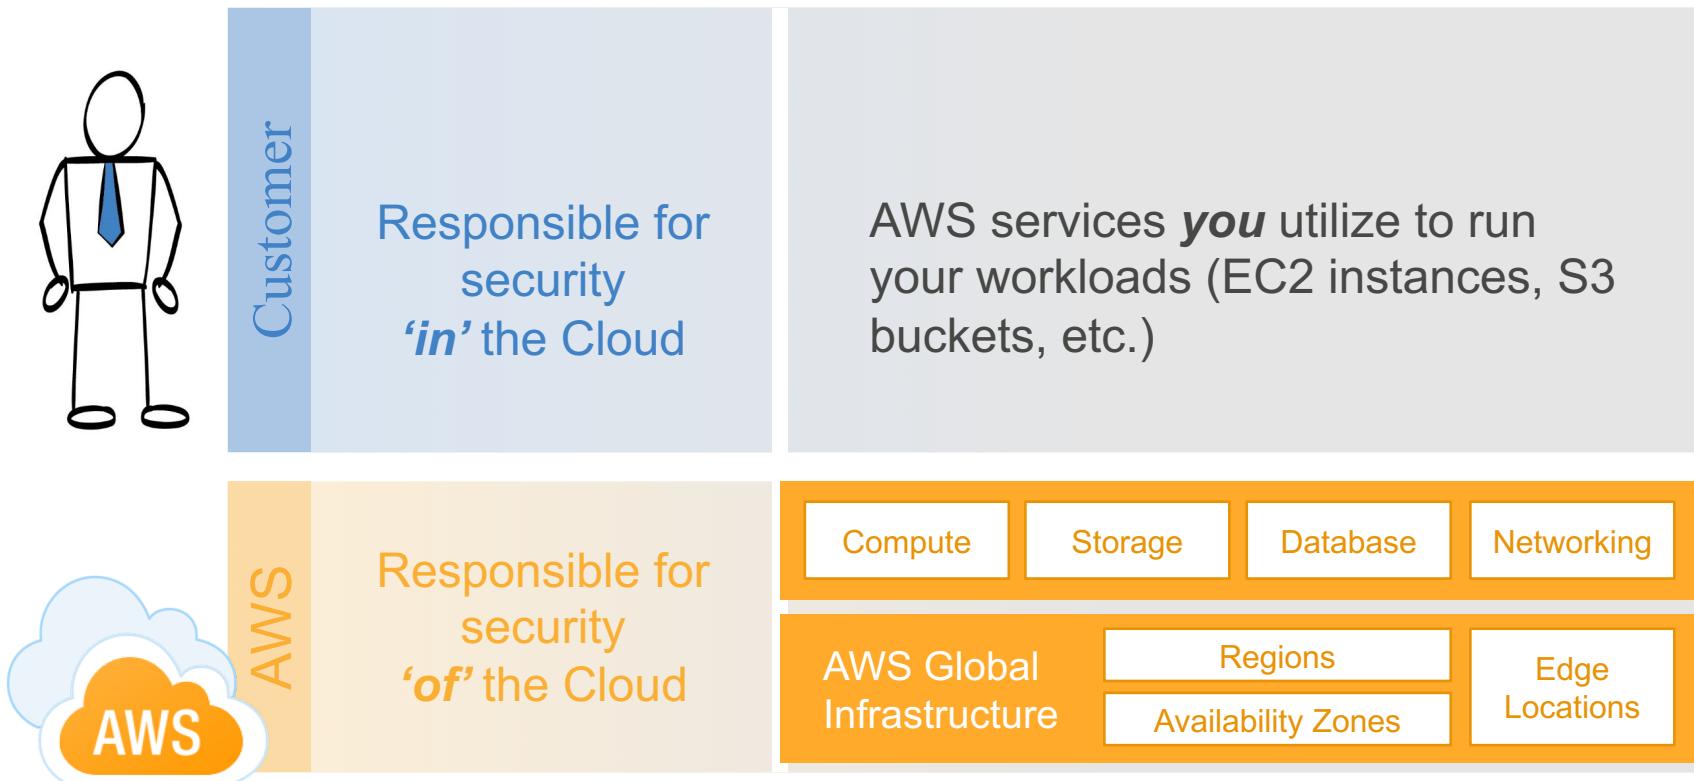
Responsible for  
security  
**'of'** the Cloud

Policies and mechanisms **AWS**  
uses to protect the cloud itself

# AWS Shared Responsibility Model



# AWS Shared Responsibility Model



# AWS Shared Responsibility Model

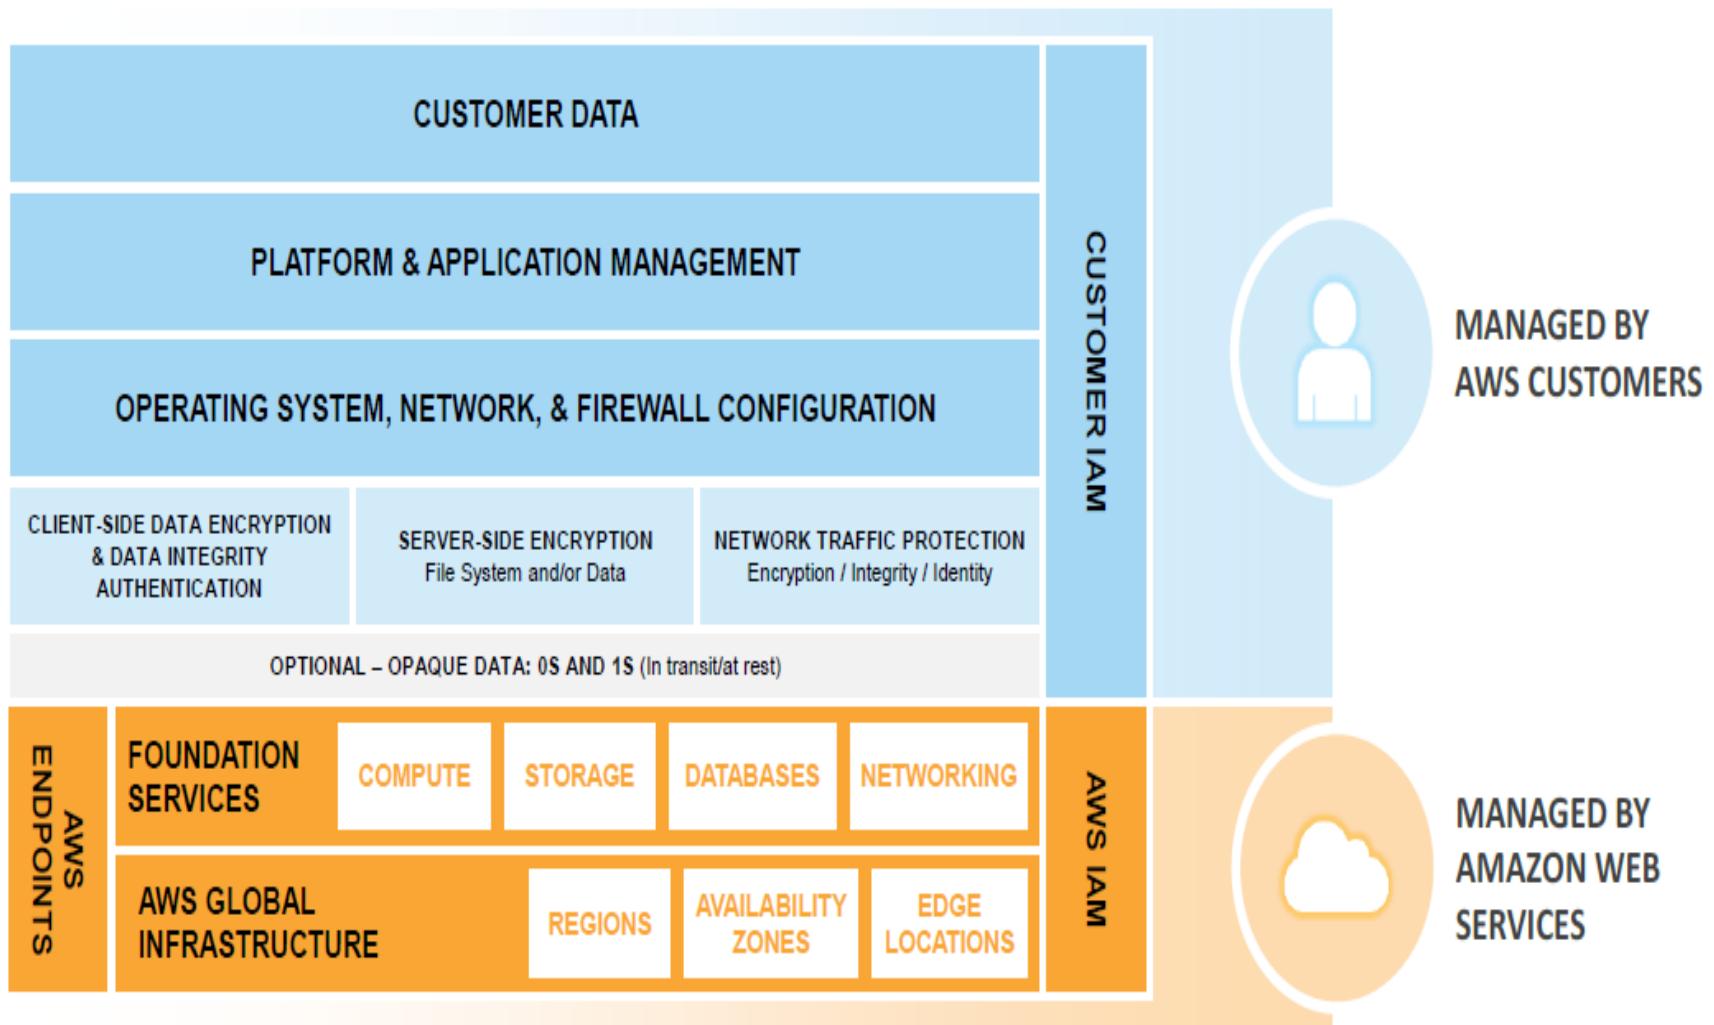


Figure 1: Shared Responsibility Model for Infrastructure Services

# Manual Configuration Challenges

---

Creating and configuring AWS services and resources through a management console is a **manual** process

What are the challenges and concerns for a manual process?

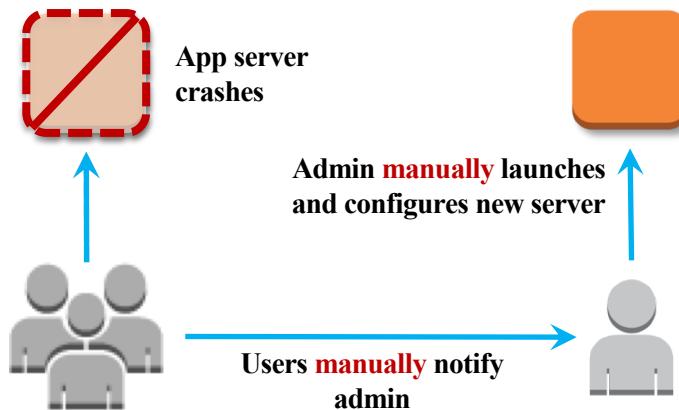
- Reliability
- Reproducibility
  - DEV
  - TEST
  - PROD
- Documentation

# Best Practice: Automate Your Environment

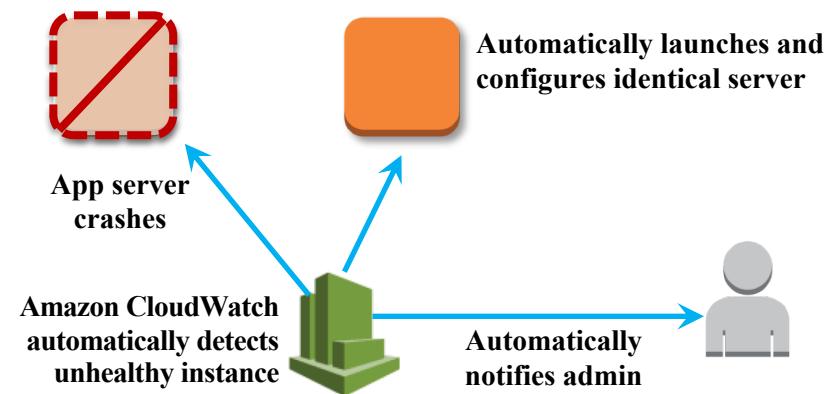
*Where possible, automate the provisioning, termination, and configuration of resources.*

Improve your system's **stability** and **consistency**, as well as the **efficiency** of your organization, by removing manual processes.

## Anti-pattern



## Best practice



# Best Practice: Use Disposable Resources

---

*Take advantage of the dynamically provisioned nature of cloud computing.*

Think of servers and other components as **temporary resources**.

## Anti-pattern

- Over time, different servers end up in different configurations
- Resources run when not needed
- Hardcoded IP addresses prevent flexibility
- Difficult/inconvenient to test new updates on hardware that's in use

## Best practice

- Automate deployment of new resources with identical configurations
- Terminate resources not in use
- Switch to new IP addresses automatically
- Test updates on new resources, and then replace old resources with updated ones

# What Does Infrastructure As Code Mean?

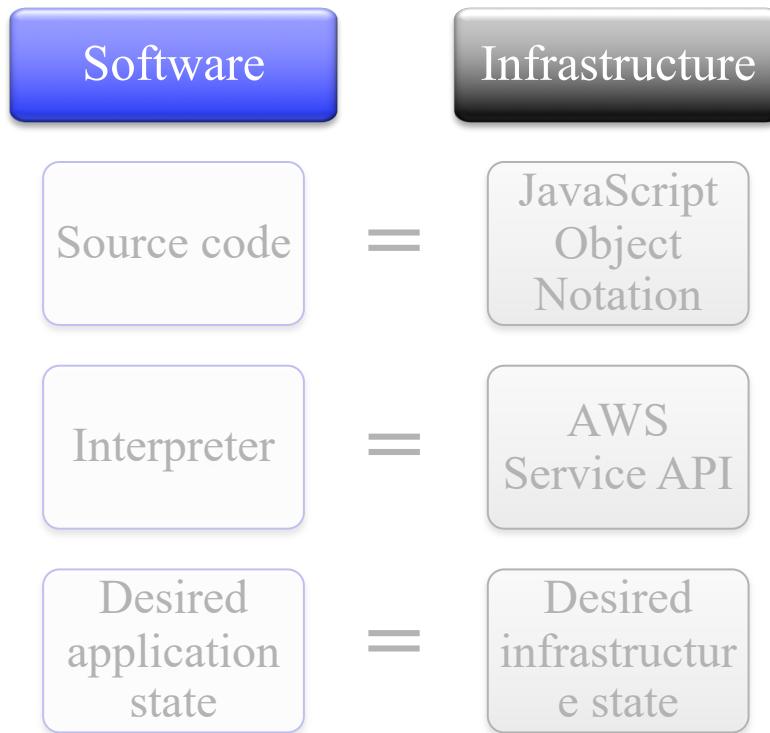
---



Techniques, practices, and tools from software development applied to creating reusable, maintainable, extensible and testable infrastructure.

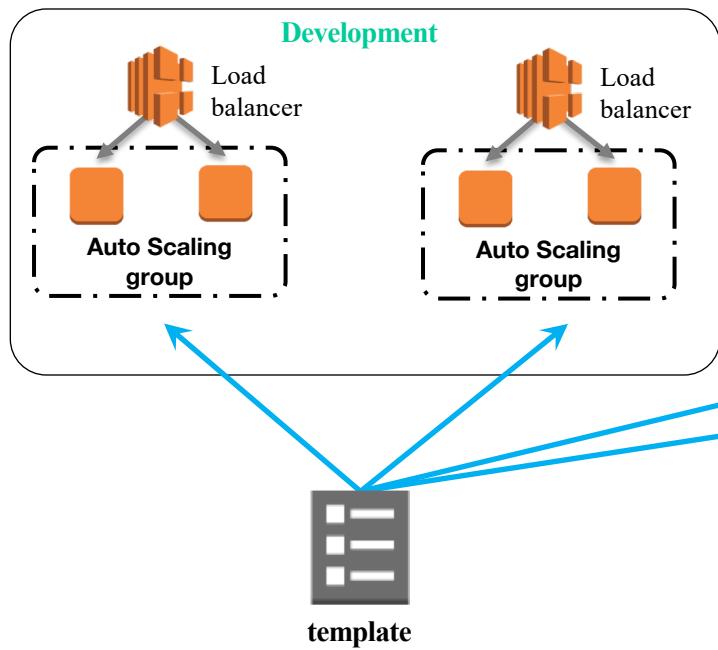
# Build And Operate Your Infrastructure Like Software

---

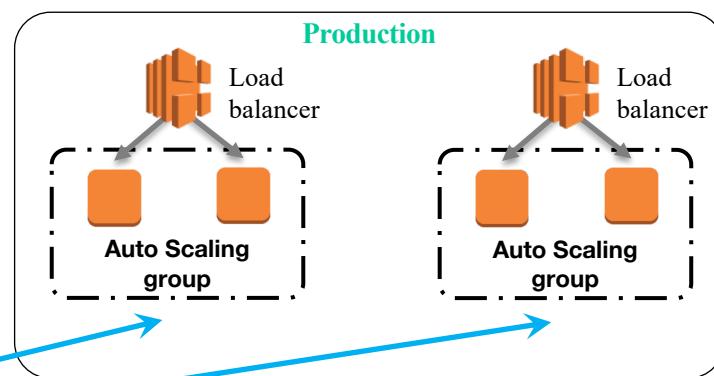


# Benefits Of Treating Infrastructure As Code

## Repeatability

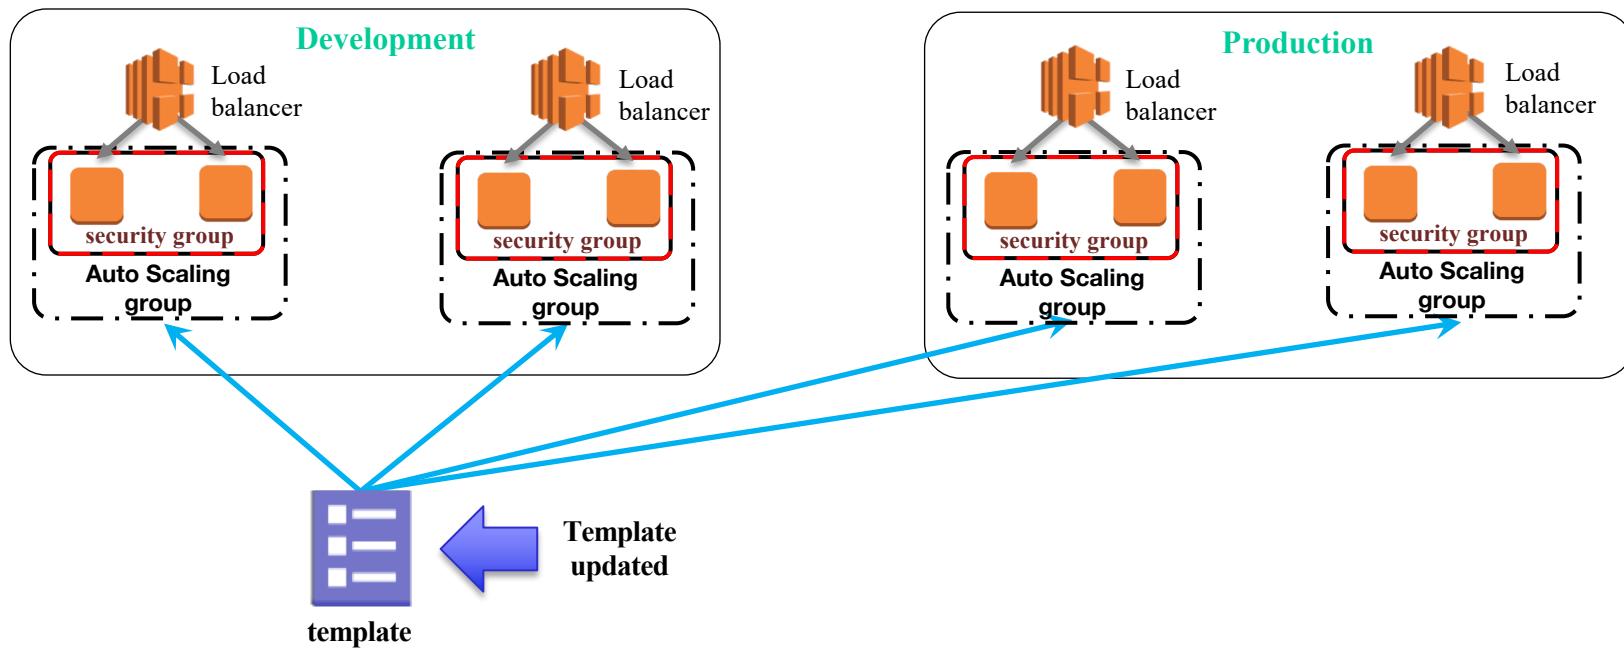


## Reusability



# Benefits Of Treating Infrastructure As Code

## Maintainability, Consistency, and Parallelization



# Infrastructure as Code : Idempotent

---

- Many popular tools used for Infrastructure as Code :
  - Ansible
  - Chef
  - Terraform
  - Puppet
- Ansible documentation – Idempotency : An operation is idempotent if the result of performing it once is exactly the same as the result of performing it repeatedly without any intervening actions.
- Compares desired state with current state to assess if changes required
- Is CloudFormation Idempotent ? Initially was not but now maybe. Partially ?

# CloudFormation: Infrastructure As Code

AWS CloudFormation allows you to launch, configure, and connect AWS resources with JavaScript Object Notation (JSON) and YAML-formatted templates.

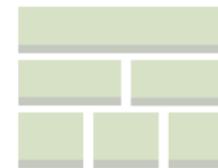
Template



AWS CloudFormation Engine



Stack



- JSON-formatted file describing the resources to be created
- Treat it as source code: put it in your repository

- AWS service component
- Interprets AWS CloudFormation template into stacks of AWS resources

- A collection of resources created by AWS CloudFormation
- Tracked and reviewable in the AWS Management Console

# Cloudformer

---

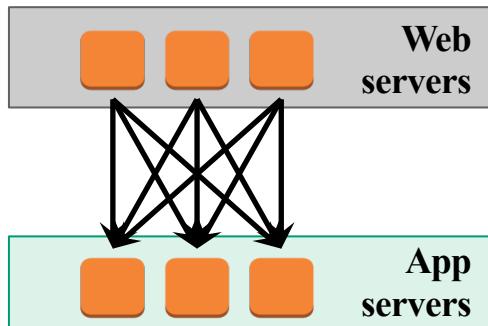
- CloudFormer is a template creation beta tool that creates an AWS CloudFormation template from existing AWS resources in your account. You select any supported AWS resources that are running in your account, and CloudFormer creates a template in an Amazon S3 bucket.
- Use CloudFormer to produce templates that you can use as a starting point. Not all AWS resources or resource properties are supported.

# Loosely Couple Your Components

*Design architectures with independent components.*

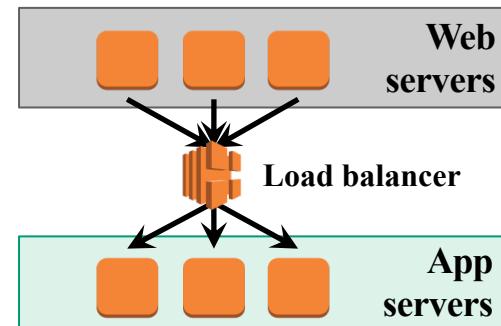
Reduce interdependencies so that the change or failure of one component does not affect other components.

## Anti-pattern



Web servers **tightly coupled** to app servers

## Best practice

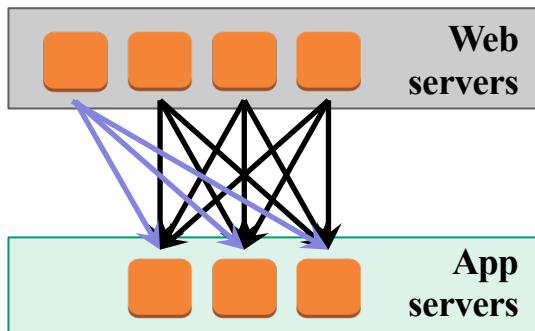


Decoupled with a load balancer

# Decoupling

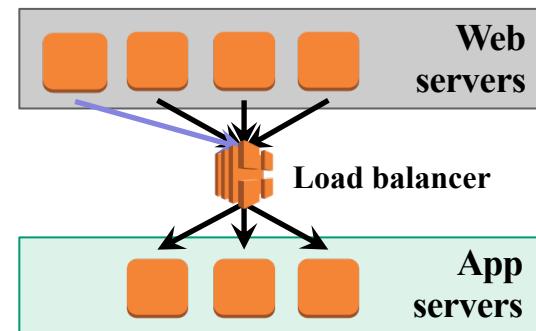
The more loosely your system is coupled...

**Tightly coupled**



the more easily  
it scales.

**Loosely coupled**

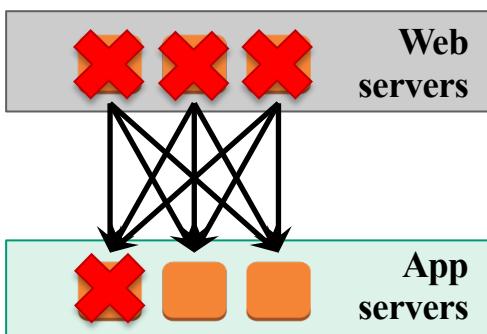


# Decoupling

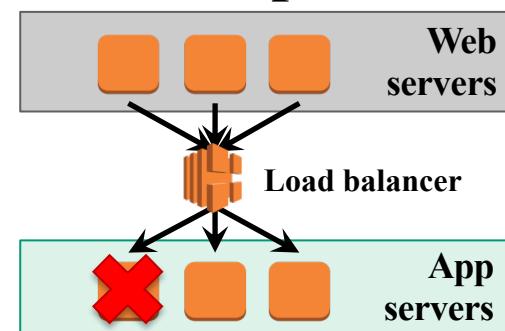
The more loosely your system is coupled...

the more easily it scales.  
the more fault-tolerant it  
can be.

**Tightly  
coupled**

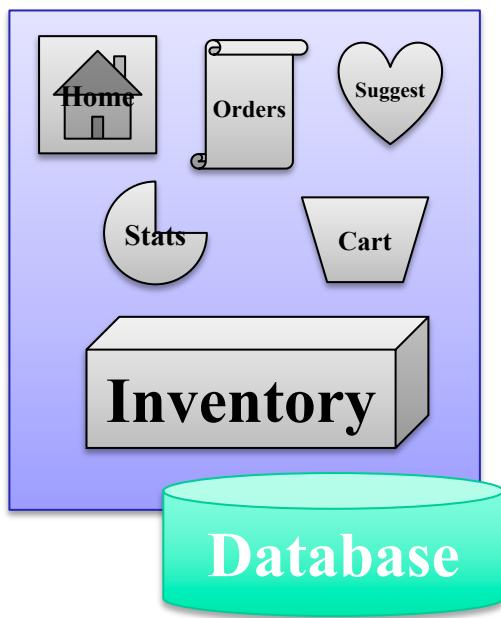


**Loosely  
coupled**

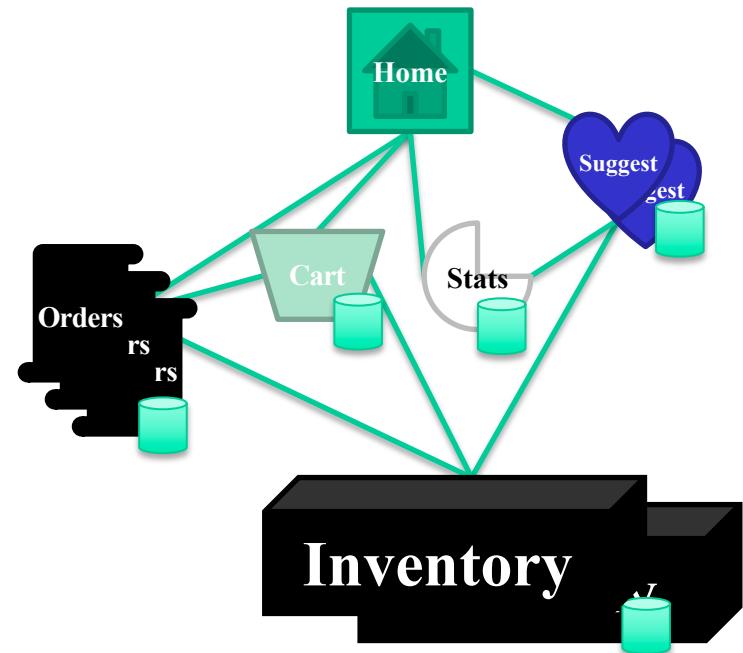


# Comparing Architectural Styles

Traditional app architectures  
are **monolithic**:



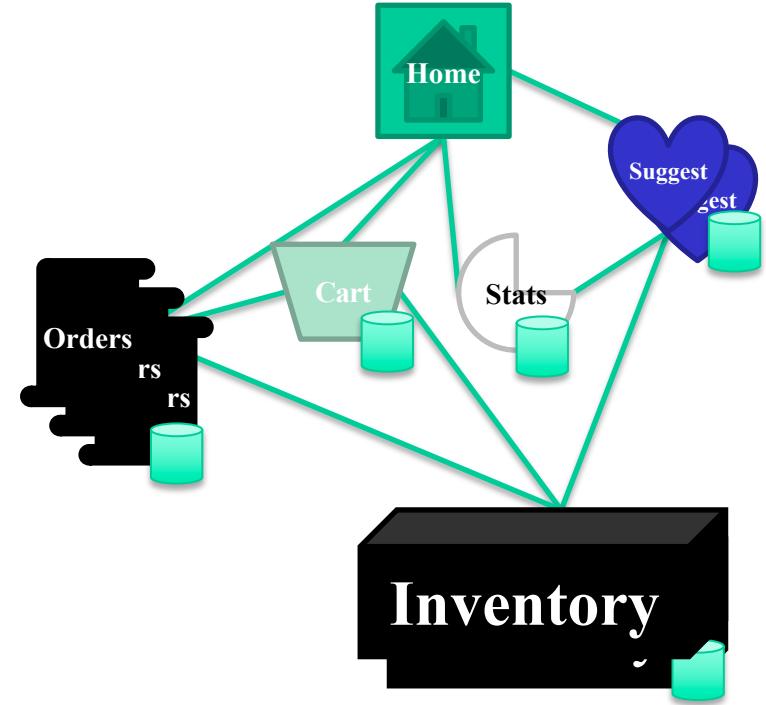
Microservice-based architectures  
are **loosely coupled**:



# Microservices

---

- Split features into individual components
- Have smaller parts to iterate on
- Have a reduced test surface area
- Benefit from a lower risk of change
- Use individual horizontally scalable parts



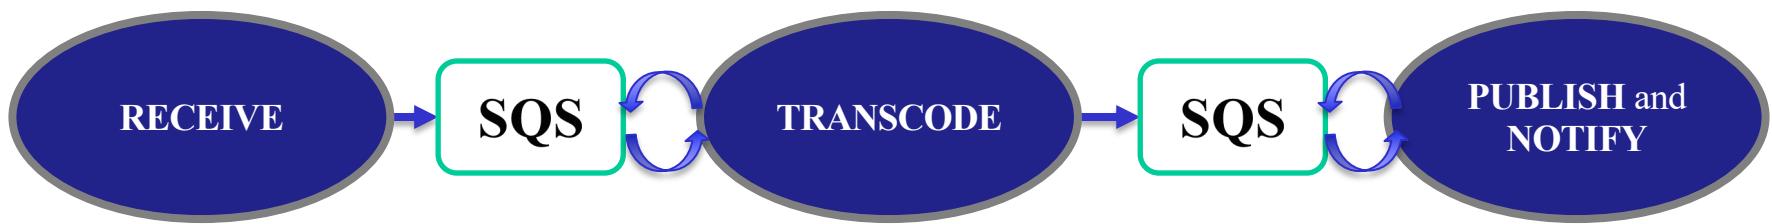
# Tightly Coupled Systems

---



# Loosely Coupled Systems

---



# Amazon Simple Queue Service (SQS)



Amazon SQS is a fully managed message queueing service. Transmit any volume of messages at any level of throughput without losing messages or requiring other services to be always available.

Messages



- Generated by one component to be consumed by another.
- Can contain 256 KB of text in any format.

Amazon SQS



- Ensures delivery of each message at least once.
- Supports multiple readers and writers on the same queue.
- Does not guarantee first in, first out.

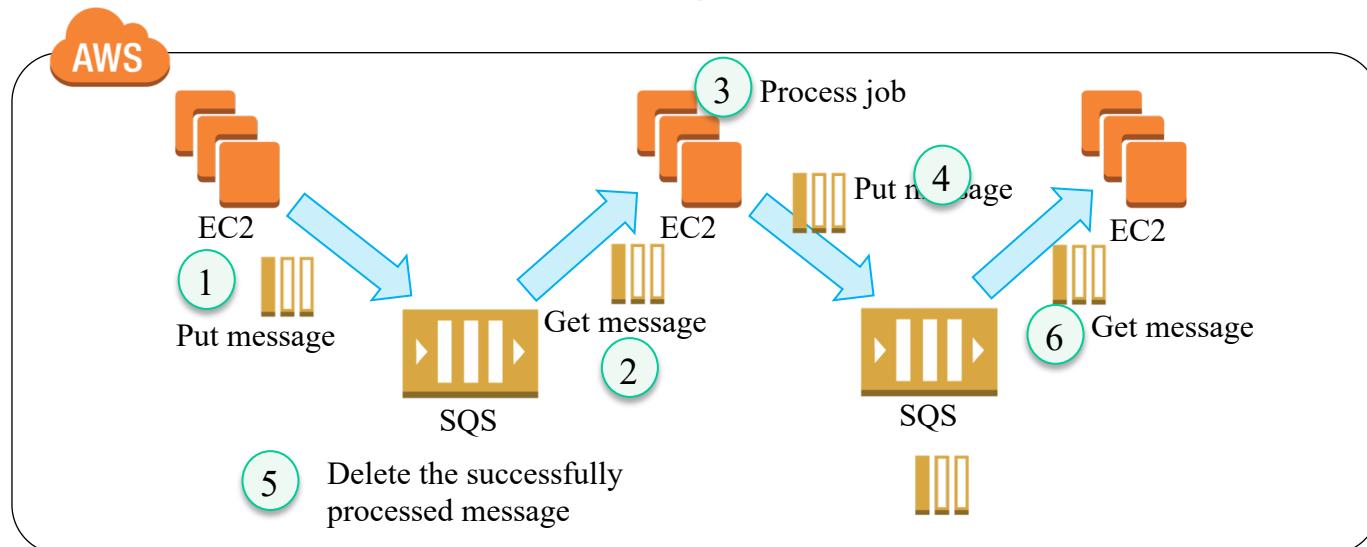
Queues



- Repository for messages awaiting processing.
- Acts as a buffer between the components which produce and receive data.

# Loose Coupling With Amazon SQS

The queuing chain pattern enables asynchronous processing:



# Amazon Simple Notification Service (SNS)

---



Amazon SNS enables you to set up, operate, and send notifications to subscribing services other applications.

- Messages published to topic
- Topic subscribers receive message

## **Subscriber types:**

- Email (plain or JSON)
- HTTP/HTTPS
- Short Message Service (SMS) clients (USA only)
- Amazon SQS queues
- Mobile push messaging
- AWS Lambda Function

# Characteristics Of Amazon SNS

---



- Single published message
- Order is not guaranteed
- No recall
- HTTP/HTTPS retry
- 256 KB max per message

# How Is Amazon SNS Different From Amazon SQS?

Amazon SQS and Amazon SNS are both messaging services within AWS.

	 <b>Amazon SNS</b>	 <b>Amazon SQS</b>
Message persistence	No	Yes
Delivery mechanism	Push (Passive)	Poll (Active)
Producer/consumer	Publish/subscribe	Send/receive

# Serverless Computing

---

- **Serverless computing** is a cloud-computing execution model in which the cloud provider acts as the server, dynamically managing the allocation of machine resources. Pricing is based on the actual amount of resources consumed by an application, rather than on pre-purchased units of capacity. It is a form of utility computing.
- The name "serverless computing" is used because the server management and capacity planning decisions are completely hidden from the developer or operator.

# AWS Lambda

---

- AWS Lambda is a compute service that lets you run code without provisioning or managing servers. Lambda executes your code only when needed and scales automatically, from a few requests per day to thousands per second. You pay only for the compute time you consume - there is no charge when your code is not running.
- AWS Lambda runs your code on a high-availability compute infrastructure and performs all of the administration of the compute resources, including server and operating system maintenance, capacity provisioning and automatic scaling, code monitoring and logging. All you need to do is supply your code in one of the languages that AWS Lambda supports (currently Node.js, Java, C# and Python).
- You can use AWS Lambda to run your code in response to events, such as changes to data in an Amazon S3 bucket or an Amazon DynamoDB table; to run your code in response to HTTP requests using Amazon API Gateway; or invoke your code using API calls made using AWS SDKs. With these capabilities, you can use Lambda to easily build data processing triggers for AWS services like Amazon S3 and Amazon DynamoDB

# Use AWS Lambda To Decouple Your Infrastructure

---



AWS Lambda is a great solution for processing data with high availability and a limited cost footprint.

AWS Lambda allows you to further decouple your infrastructure by replacing traditional servers with simple microprocesses.

# Serverless Computing With AWS Lambda

---



AWS Lambda starts code within milliseconds of an event such as:

- An image upload
- In-app activity
- A website click
- Output from a connected device

**Consider AWS Lambda if:**

- You're using entire instances to run simple functions or processing applications
- You don't want to worry about HA, scaling, deployment, or management

# Triggers For AWS Lambda Functions

---



Amazon  
DynamoDB



Amazon  
S3



Amazon  
SNS



Amazon  
Kinesis



Amazon  
SES



AWS  
Config



Scheduled  
Events



Amazon  
Cognito



AWS SDKs  
via Amazon  
API Gateway



Amazon  
CloudWatch



AWS  
CodeCommit



AWS  
CloudFormation

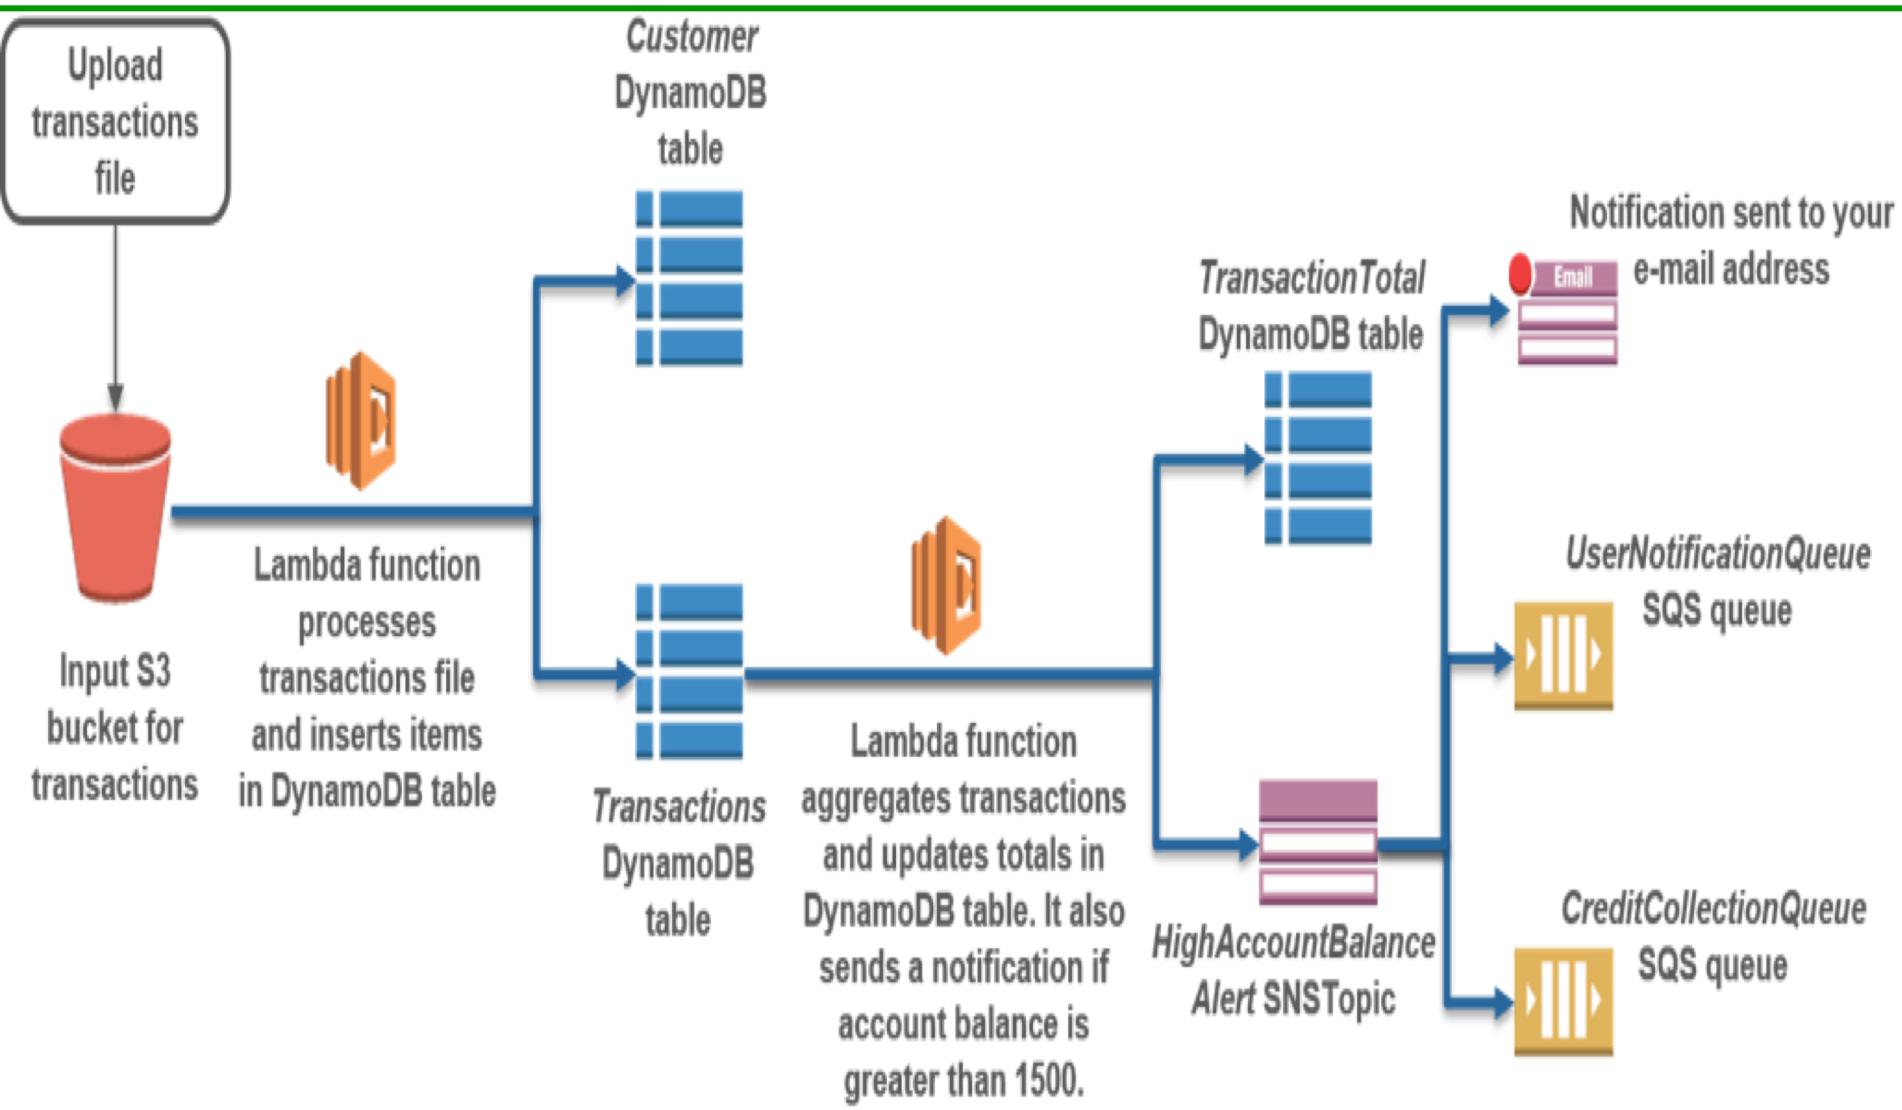


Amazon Echo:  
Alexa Skills



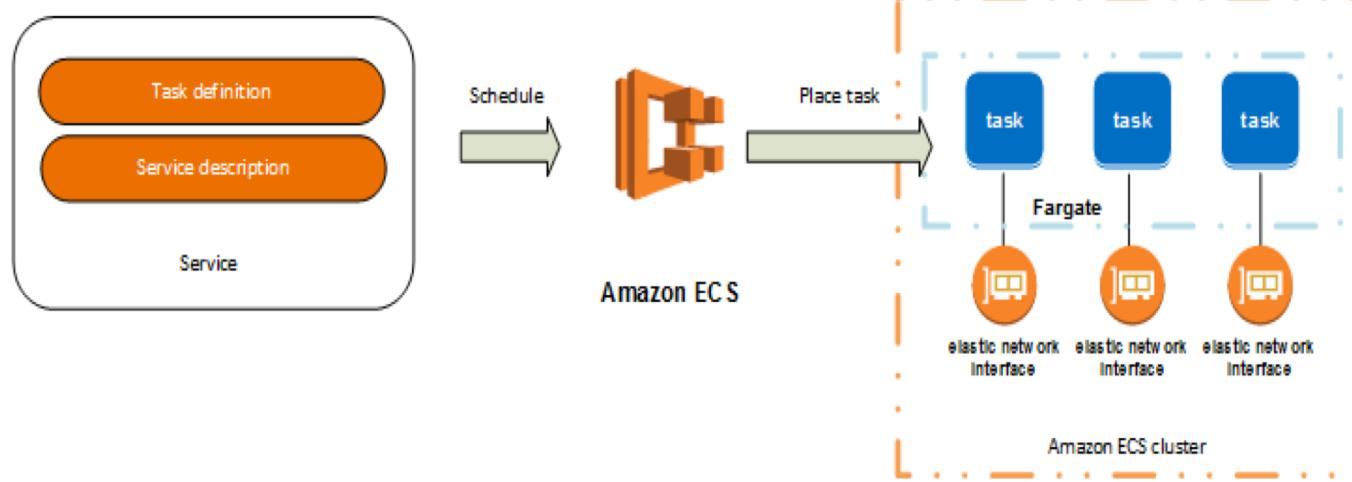
HTTPS via API  
Gateway

# AWS Lambda



# Amazon ECS

- Amazon Elastic Container Service (Amazon ECS) is a highly scalable, fast, container management service that makes it easy to run, stop, and manage Docker containers on a cluster. You can host your cluster on a serverless infrastructure that is managed by Amazon ECS by launching your services or tasks using the Fargate launch type. For more control you can host your tasks on a cluster of Amazon Elastic Compute Cloud (Amazon EC2) instances that you manage by using the EC2 launch type.



# Amazon CloudWatch



Amazon  
CloudWatch

- A **monitoring service** for AWS cloud resources and the applications you run on AWS
- **Visibility into** resource utilization, operational performance, and overall demand patterns
- **Custom application-specific** metrics of your own
- **Accessible** via AWS Management Console, APIs, SDK, or CLI

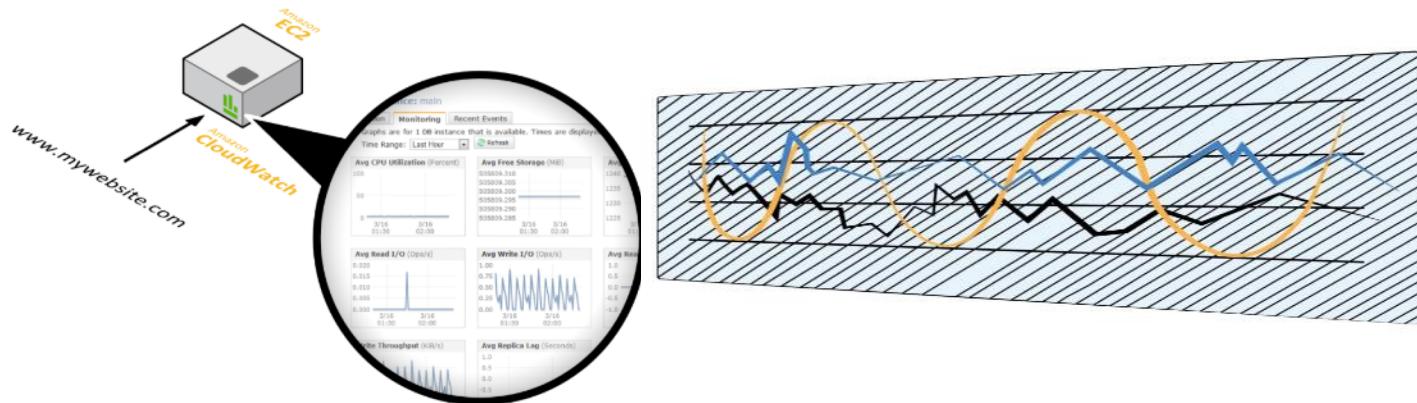
# Amazon CloudWatch Facts



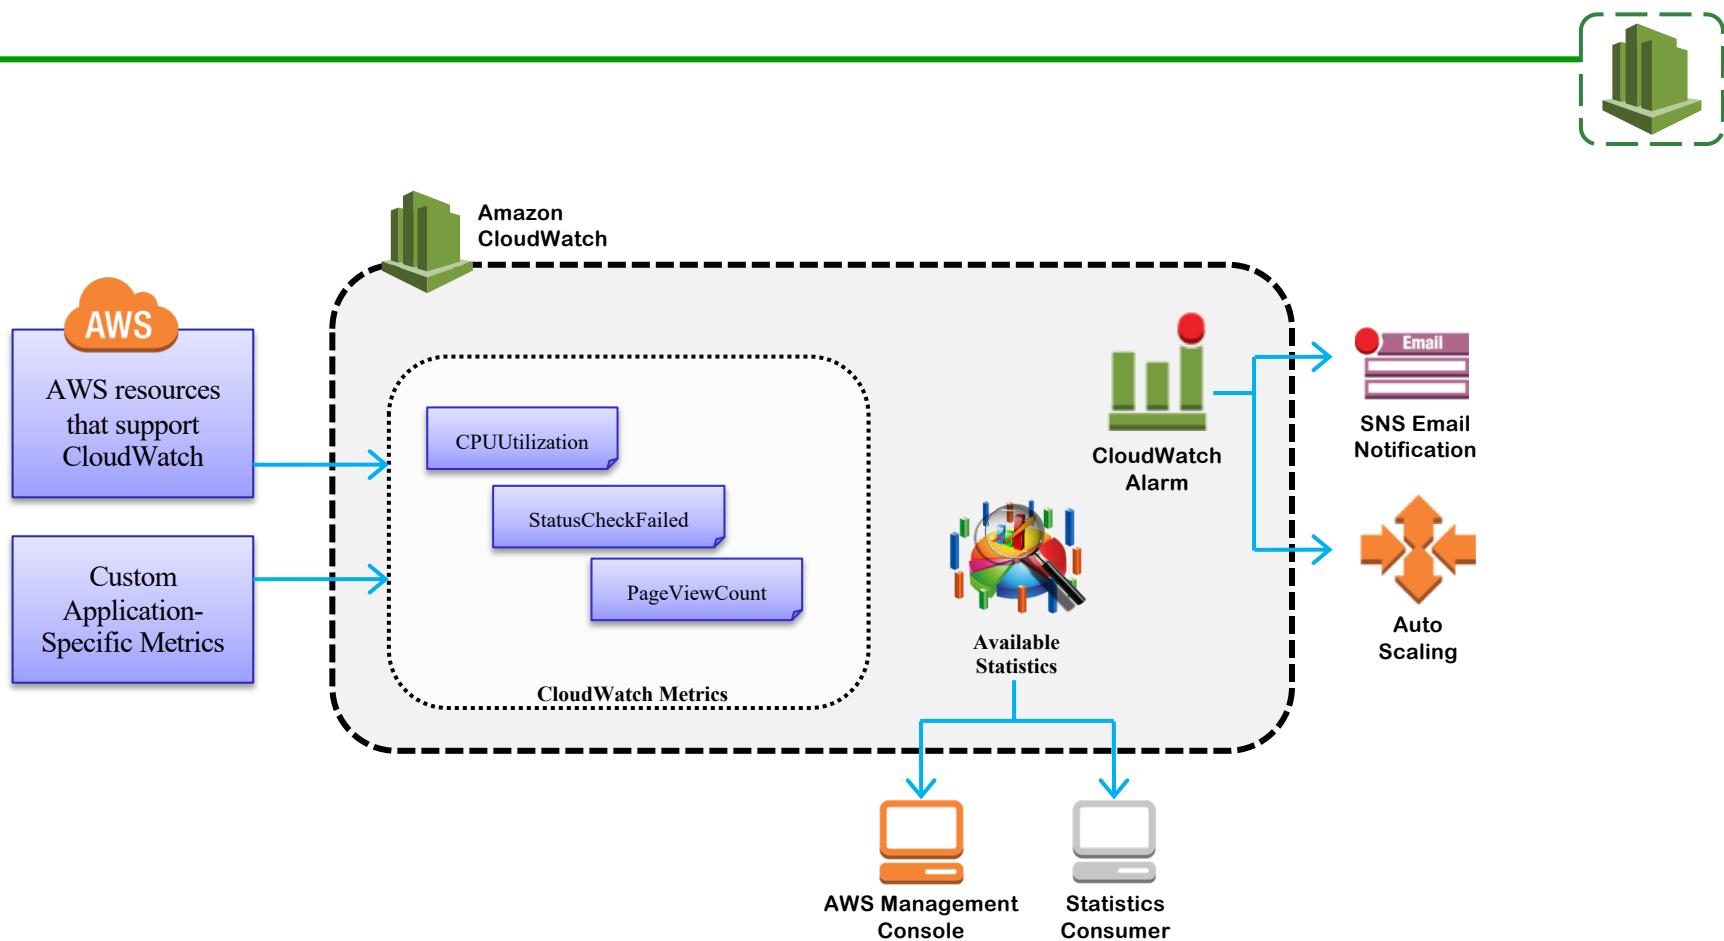
- Monitor other AWS resources

*View graphics and statistics*

- Set Alarms



# Amazon CloudWatch Architecture



# CloudWatch Metrics Examples



CloudWatch Metrics by Category

Your CloudWatch metric summary has loaded. Total metrics: 97

EBS Metrics: 24      EC2 Metrics: 38      S3 Metrics: 18      SNS Metrics: 3      SQS Metrics: 14  
Per-Volume Metrics: 24      Per-Instance Metrics: 38      Storage Metrics: 18      Topic Metrics: 3      Queue Metrics: 14

Browse Metrics ▾  X

Showing all results (100) for cpu. For more results expand your search to All EC2 Metrics.  
[Select All](#) | [Clear](#)

EC2 > Per-Instance Metrics

InstanceId	InstanceName	Metric Name
<input checked="" type="checkbox"/> i-2be9e804		CPUUtilization
<input type="checkbox"/> i-2ce7e603		CPUUtilization
<input type="checkbox"/> i-5b210d23		CPUUtilization
<input type="checkbox"/> i-8da34f60		CPUCreditBalance

Title: CPUUtilization  Average

A line chart titled "CPUUtilization" showing CPU utilization over a 5-minute period. The Y-axis ranges from 0 to 0.06 with increments of 0.02. The X-axis shows time from 12:00 to 23:00. The chart displays a high-frequency, low-amplitude oscillation between approximately 0.01 and 0.06, indicating fluctuating CPU usage.

Actions: Add to Dashboard, Copy URL, Create Alarm

Time Range: Relative, Absolute, UTC (GMT)  
From: 12 hours ago, To: 0 hours ago  
Zoom: 1h | 3h | 6h | 12h | 1d | 3d | 1w | 2w

Left Y-axis

# CloudWatch Custom Metrics

---

- You can publish your own metrics to CloudWatch using the AWS CLI or an API. You can view statistical graphs of your published metrics with the AWS Management Console.
- CloudWatch stores data about a metric as a series of data points. Each data point has an associated time stamp.